



A Gentle Introduction to Gaussian Processes

John Joseph Valletta

University of Exeter, Penryn Campus, UK

Internal Maths Seminar: 20th October 2015



- Motivation
- The Gaussian Distribution
- Gaussian Processes
- Gaussian Process Regression - A Toy Example
- Gaussian Process Regression - CO₂ Concentrations
- Modelling Gene Expression Time-Series

- Motivation
- The Gaussian Distribution
- Gaussian Processes
- Gaussian Process Regression - A Toy Example
- Gaussian Process Regression - CO₂ Concentrations
- Modelling Gene Expression Time-Series


- Motivation
- The Gaussian Distribution
- Gaussian Processes
- Gaussian Process Regression - A Toy Example
- Gaussian Process Regression - CO₂ Concentrations
- Modelling Gene Expression Time-Series

- Motivation
- The Gaussian Distribution
- Gaussian Processes
- Gaussian Process Regression - A Toy Example
- Gaussian Process Regression - CO₂ Concentrations
- Modelling Gene Expression Time-Series


- Motivation
- The Gaussian Distribution
- Gaussian Processes
- Gaussian Process Regression - A Toy Example
- Gaussian Process Regression - CO₂ Concentrations
- Modelling Gene Expression Time-Series

- Motivation
- The Gaussian Distribution
- Gaussian Processes
- Gaussian Process Regression - A Toy Example
- Gaussian Process Regression - CO₂ Concentrations
- Modelling Gene Expression Time-Series


Motivation



Motivation



Motivation




The Data Modelling Task


- **data:** $\mathbf{x} = \{x_1, \dots, x_N\}$, $\mathbf{y} = \{y_1, \dots, y_N\}$
- **model:** $y = f(x) + \epsilon$
- **predictions:** $y^* = f(x^*)$

The Data Modelling Task

- **data:** $\mathbf{x} = \{x_1, \dots, x_N\}$, $\mathbf{y} = \{y_1, \dots, y_N\}$
- **model:** $y = f(x) + \epsilon$
- **predictions:** $y^* = f(x^*)$




The Gaussian Distribution



$$\mathcal{N}(\mu, \sigma^2)$$


A draw from this distribution is a 1D vector
e.g $x = [0.2]$



$$\mathcal{N}(\mu, \Sigma)$$


A draw from this distribution is a 2D vector
e.g $\mathbf{x} = \begin{bmatrix} 0.3 \\ -0.4 \end{bmatrix}$

The Covariance Matrix



Isotropic


$$\Sigma = \begin{pmatrix} 1 & 0 \\ 0 & 1 \end{pmatrix}$$



Diagonal


$$\Sigma = \begin{pmatrix} 1 & 0 \\ 0 & 0.5 \end{pmatrix}$$

The Covariance Matrix



General Form

$$\Sigma = \begin{pmatrix} 1 & 0.5 \\ 0.5 & 1 \end{pmatrix}$$



General Form

$$\Sigma = \begin{pmatrix} 1 & -0.5 \\ -0.5 & 1 \end{pmatrix}$$


Sampling from a Multivariate Gaussian

What does a *single* sample from a 100 dimensional Gaussian look like?

$$\mathbf{x} = \left[\begin{array}{c} 0.2 \\ \vdots \\ \vdots \\ \vdots \\ \vdots \\ \vdots \\ \vdots \\ \vdots \\ -0.6 \end{array} \right] \Bigg\} = 100 \Rightarrow$$

Sampling from a Multivariate Gaussian

What does a *single* sample from a 100 dimensional Gaussian look like?



Gaussian Process in a Nutshell

Recall: What we are after is $y = f(x)$

Trick: Think about a function as an infinitely-long vector


$$f(x) = \left[\begin{array}{c} f_1 \\ \vdots \\ \vdots \\ \vdots \\ \vdots \\ \vdots \\ \vdots \\ \vdots \\ f_\infty \end{array} \right] \Bigg\} = \infty \Rightarrow$$

Gaussian Process in a Nutshell

Recall: What we are after is $y = f(x)$

Trick: Think about a function as an infinitely-long vector

$$f(x) = \left[\begin{array}{c} f_1 \\ \vdots \\ \vdots \\ \vdots \\ \vdots \\ \vdots \\ \vdots \\ \vdots \\ f_\infty \end{array} \right] \Bigg\} = \infty \Rightarrow$$




Gaussian Process in a Nutshell

Recall: What we are after is $y = f(x)$

Trick: Think about a function as an infinitely-long vector

$$f(x) = \left[\begin{array}{c} f_1 \\ \vdots \\ \vdots \\ \vdots \\ \vdots \\ \vdots \\ \vdots \\ \vdots \\ f_\infty \end{array} \right] = \infty \Rightarrow$$



Computational Madness: Ask only for the properties of the function at a *finite* number of points

Gaussian Process in a Nutshell

3 dimensional Gaussian

Mean vector

$$\mu = \begin{pmatrix} \mu_1 \\ \mu_2 \\ \mu_3 \end{pmatrix}$$

Covariance matrix

$$\Sigma = \begin{pmatrix} \sigma_1^2 & \sigma_{12} & \sigma_{13} \\ \sigma_{21} & \sigma_2^2 & \sigma_{23} \\ \sigma_{31} & \sigma_{32} & \sigma_3^2 \end{pmatrix}$$

∞ dimensional Gaussian

Mean *function*

$$\mu = \begin{pmatrix} \mu_1 \\ \vdots \\ \mu_\infty \end{pmatrix} = m(\mathbf{x})$$

Covariance *function*

$$\Sigma = \begin{pmatrix} \sigma_1^2 & \sigma_{12} & \dots & \sigma_{1\infty} \\ \sigma_{21} & \sigma_2^2 & \dots & \vdots \\ \vdots & \vdots & \ddots & \vdots \\ \vdots & \vdots & \dots & \sigma_\infty^2 \end{pmatrix} = k(\mathbf{x}, \mathbf{x}')$$

Definition

A Gaussian Process (GP) is an infinite collection of random variables, any finite number of which have a joint normal distribution.

Essentially an infinite dimension multivariate Gaussian distribution, characterised by a mean function $m(\mathbf{x})$ and a covariance function $k(\mathbf{x}, \mathbf{x}')$

Rationale

Instead of inferring the parameters of a fixed model structure $(\beta_0, \beta_1, \dots)$, with GPs we model the *correlation* between inputs. That is, inputs \mathbf{x} that are close/similar to each other are likely to give rise to a similar output $f(\mathbf{x})$

$$\begin{aligned}f(\mathbf{x}) &\sim \text{GP}(m(\mathbf{x}), k(\mathbf{x}, \mathbf{x}')) \\m(\mathbf{x}) &= \text{E}[f(\mathbf{x})] \\k(\mathbf{x}, \mathbf{x}') &= \text{E}[(f(\mathbf{x}) - m(\mathbf{x}))(f(\mathbf{x}') - m(\mathbf{x}'))]\end{aligned}$$

Covariance Function

- Vital ingredient in Gaussian Process¹
 - Encodes our assumptions about the function we wish to model (smooth, stationary, etc.)
 - Quantifies the *similarity* between two data points; crucial for predicting a test point \mathbf{x}^*
 - Needs to satisfy a set of mathematical conditions (beyond the scope of this intro)
 - A very popular choice is the Squared Exponential:
(also known as RBF, Gaussian and Exponentiated Quadratic Kernel Function)

$$k(\mathbf{x}, \mathbf{x}') = \alpha \exp \left(-\frac{\|\mathbf{x} - \mathbf{x}'\|^2}{2l^2} \right)$$

e.g if we set $\alpha = 1$ and $l = 1$ then:

$$k(0, 0) = e^0 = 1, \quad k(0, 1) = e^{-\frac{1}{2}} = 0.6, \quad k(0, 2) = e^{-2} = 0.14$$

¹without much loss of generality we can assume that $m(\mathbf{x}) \equiv 0$

Covariance Function

- Vital ingredient in Gaussian Process¹
- Encodes our assumptions about the function we wish to model (smooth, stationary, etc.)
- Quantifies the *similarity* between two data points; crucial for predicting a test point \mathbf{x}^*
- Needs to satisfy a set of mathematical conditions (beyond the scope of this intro)
- A very popular choice is the Squared Exponential:
(also known as RBF, Gaussian and Exponentiated Quadratic Kernel Function)

$$k(\mathbf{x}, \mathbf{x}') = \alpha \exp \left(-\frac{\|\mathbf{x} - \mathbf{x}'\|^2}{2l^2} \right)$$

e.g if we set $\alpha = 1$ and $l = 1$ then:

$$k(0, 0) = e^0 = 1, \quad k(0, 1) = e^{-\frac{1}{2}} = 0.6, \quad k(0, 2) = e^{-2} = 0.14$$

¹without much loss of generality we can assume that $m(\mathbf{x}) \equiv 0$

Covariance Function

- Vital ingredient in Gaussian Process¹
- Encodes our assumptions about the function we wish to model (smooth, stationary, etc.)
- Quantifies the *similarity* between two data points; crucial for predicting a test point \mathbf{x}^*
- Needs to satisfy a set of mathematical conditions (beyond the scope of this intro)
- A very popular choice is the Squared Exponential:
(also known as RBF, Gaussian and Exponentiated Quadratic Kernel Function)

$$k(\mathbf{x}, \mathbf{x}') = \alpha \exp \left(-\frac{\|\mathbf{x} - \mathbf{x}'\|^2}{2l^2} \right)$$

e.g if we set $\alpha = 1$ and $l = 1$ then:

$$k(0, 0) = e^0 = 1, \quad k(0, 1) = e^{-\frac{1}{2}} = 0.6, \quad k(0, 2) = e^{-2} = 0.14$$

¹without much loss of generality we can assume that $m(\mathbf{x}) \equiv 0$

Covariance Function

- Vital ingredient in Gaussian Process¹
- Encodes our assumptions about the function we wish to model (smooth, stationary, etc.)
- Quantifies the *similarity* between two data points; crucial for predicting a test point \mathbf{x}^*
- Needs to satisfy a set of mathematical conditions (beyond the scope of this intro)
- A very popular choice is the Squared Exponential:
(also known as RBF, Gaussian and Exponentiated Quadratic Kernel Function)

$$k(\mathbf{x}, \mathbf{x}') = \alpha \exp \left(-\frac{\|\mathbf{x} - \mathbf{x}'\|^2}{2l^2} \right)$$

e.g if we set $\alpha = 1$ and $l = 1$ then:

$$k(0, 0) = e^0 = 1, \quad k(0, 1) = e^{-\frac{1}{2}} = 0.6, \quad k(0, 2) = e^{-2} = 0.14$$

¹without much loss of generality we can assume that $m(\mathbf{x}) \equiv 0$

Covariance Function

- Vital ingredient in Gaussian Process¹
- Encodes our assumptions about the function we wish to model (smooth, stationary, etc.)
- Quantifies the *similarity* between two data points; crucial for predicting a test point \mathbf{x}^*
- Needs to satisfy a set of mathematical conditions (beyond the scope of this intro)
- A very popular choice is the Squared Exponential:
(also known as RBF, Gaussian and Exponentiated Quadratic Kernel Function)

$$k(\mathbf{x}, \mathbf{x}') = \alpha \exp \left(-\frac{\|\mathbf{x} - \mathbf{x}'\|^2}{2l^2} \right)$$

e.g if we set $\alpha = 1$ and $l = 1$ then:

$$k(0, 0) = e^0 = 1, \quad k(0, 1) = e^{-\frac{1}{2}} = 0.6, \quad k(0, 2) = e^{-2} = 0.14$$

¹without much loss of generality we can assume that $m(\mathbf{x}) \equiv 0$

Gaussian Process Regression

- Shift the problem from inferring model parameters to choosing a covariance function and its (hyper)parameters
- Choose $m(\mathbf{x})$ and $k(\mathbf{x}, \mathbf{x}')$ that reflect some prior belief
- This defines a *prior* on the function class itself; it is a prior on a *function* and not parameters of some fixed model structure
- Under a Bayesian framework this prior is “reshaped” by the observed data to obtain a posterior distribution on the *function*

Gaussian Process Regression

- Shift the problem from inferring model parameters to choosing a covariance function and its (hyper)parameters
- Choose $m(\mathbf{x})$ and $k(\mathbf{x}, \mathbf{x}')$ that reflect some prior belief
- This defines a *prior* on the function class itself; it is a prior on a *function* and not parameters of some fixed model structure
- Under a Bayesian framework this prior is “reshaped” by the observed data to obtain a posterior distribution on the *function*


Gaussian Process Regression

- Shift the problem from inferring model parameters to choosing a covariance function and its (hyper)parameters
- Choose $m(\mathbf{x})$ and $k(\mathbf{x}, \mathbf{x}')$ that reflect some prior belief
- This defines a *prior* on the function class itself; it is a prior on a *function* and not parameters of some fixed model structure
- Under a Bayesian framework this prior is “reshaped” by the observed data to obtain a posterior distribution on the *function*


Gaussian Process Regression

- Shift the problem from inferring model parameters to choosing a covariance function and its (hyper)parameters
- Choose $m(\mathbf{x})$ and $k(\mathbf{x}, \mathbf{x}')$ that reflect some prior belief
- This defines a *prior* on the function class itself; it is a prior on a *function* and not parameters of some fixed model structure
- Under a Bayesian framework this prior is “reshaped” by the observed data to obtain a posterior distribution on the *function*


Gaussian Process Regression - A Toy Example




Gaussian Process Regression - A Toy Example




Gaussian Process Regression - A Toy Example




Gaussian Process Regression - A Toy Example




Gaussian Process Regression - A Toy Example




Gaussian Process Regression - A Toy Example



Gaussian Process Regression - A Toy Example



Gaussian Process Regression - A Toy Example



Misspecifying the Covariance Function

- The Squared Exponential covariance function was used in the previous example:

$$k(\mathbf{x}, \mathbf{x}') = \alpha \exp \left(-\frac{\|\mathbf{x} - \mathbf{x}'\|^2}{2l^2} \right)$$

- Choosing a covariance functions is akin to choosing a model structure; it dictates the class of functions that can be represented by the Gaussian Process²
- Misspecifying the covariance function and/or its (hyper)parameters has a detrimental effect on the model fit
- For e.g in the squared exponential case the lengthscale l dictates how much the function is allowed to bend

²typically a wider class of functions than in parametric models

Misspecifying the Covariance Function

- The Squared Exponential covariance function was used in the previous example:

$$k(\mathbf{x}, \mathbf{x}') = \alpha \exp \left(-\frac{\|\mathbf{x} - \mathbf{x}'\|^2}{2l^2} \right)$$

- Choosing a covariance functions is akin to choosing a model structure; it dictates the class of functions that can be represented by the Gaussian Process²
- Misspecifying the covariance function and/or its (hyper)parameters has a detrimental effect on the model fit
- For e.g in the squared exponential case the lengthscale l dictates how much the function is allowed to bend

²typically a wider class of functions than in parametric models

Misspecifying the Covariance Function

- The Squared Exponential covariance function was used in the previous example:

$$k(\mathbf{x}, \mathbf{x}') = \alpha \exp \left(-\frac{\|\mathbf{x} - \mathbf{x}'\|^2}{2l^2} \right)$$

- Choosing a covariance functions is akin to choosing a model structure; it dictates the class of functions that can be represented by the Gaussian Process²
- Misspecifying the covariance function and/or its (hyper)parameters has a detrimental effect on the model fit
- For e.g in the squared exponential case the lengthscale l dictates how much the function is allowed to bend

²typically a wider class of functions than in parametric models

Misspecifying the Covariance Function


- The Squared Exponential covariance function was used in the previous example:

$$k(\mathbf{x}, \mathbf{x}') = \alpha \exp \left(-\frac{\|\mathbf{x} - \mathbf{x}'\|^2}{2l^2} \right)$$


- Choosing a covariance functions is akin to choosing a model structure; it dictates the class of functions that can be represented by the Gaussian Process²
- Misspecifying the covariance function and/or its (hyper)parameters has a detrimental effect on the model fit
- For e.g in the squared exponential case the lengthscale l dictates how much the function is allowed to bend

²typically a wider class of functions than in parametric models


The Effect of Lengthscale on Model Fit




The Effect of Lengthscale on Model Fit




The Effect of Lengthscale on Model Fit




The Effect of Lengthscale on Model Fit




The Effect of Lengthscale on Model Fit




The Effect of Lengthscale on Model Fit



The Effect of Lengthscale on Model Fit



The Effect of Lengthscale on Model Fit



Learning the Covariance Parameters

- Let us define θ as a vector containing all (hyper)parameters
e.g $\theta = \{\alpha, l, \sigma_n^2\}$

- We choose θ that maximises the log marginal likelihood, that is:

$$\ln p(y|\mathbf{x}, \theta) = -\frac{1}{2} \mathbf{y}^T (k(\mathbf{x}, \mathbf{x}') + \sigma_n^2 I)^{-1} \mathbf{y} - \frac{1}{2} \ln |k(\mathbf{x}, \mathbf{x}') + \sigma_n^2 I| - \frac{n}{2} \ln 2\pi$$

- Cannot guarantee a global optimum; try different initial conditions
- Constraint (hyper)parameters to some sensible limits
- **Note:** Choosing the right covariance function (e.g Squared Exponential, Matérn, Rational Quadratic, etc.) is not always easy and should be treated akin to a model selection problem

Learning the Covariance Parameters

- Let us define θ as a vector containing all (hyper)parameters
e.g $\theta = \{\alpha, l, \sigma_n^2\}$
- We choose θ that maximises the log marginal likelihood, that is:

$$\ln p(y|\mathbf{x}, \theta) = -\frac{1}{2}y^T(k(\mathbf{x}, \mathbf{x}') + \sigma_n^2 I)^{-1}y - \frac{1}{2} \ln |k(\mathbf{x}, \mathbf{x}') + \sigma_n^2 I| - \frac{n}{2} \ln 2\pi$$

- Cannot guarantee a global optimum; try different initial conditions
- Constraint (hyper)parameters to some sensible limits
- **Note:** Choosing the right covariance function (e.g Squared Exponential, Matérn, Rational Quadratic, etc.) is not always easy and should be treated akin to a model selection problem

Learning the Covariance Parameters

- Let us define θ as a vector containing all (hyper)parameters
e.g $\theta = \{\alpha, l, \sigma_n^2\}$
- We choose θ that maximises the log marginal likelihood, that is:

$$\ln p(y|\mathbf{x}, \theta) = -\frac{1}{2}y^T(k(\mathbf{x}, \mathbf{x}') + \sigma_n^2 I)^{-1}y - \frac{1}{2} \ln |k(\mathbf{x}, \mathbf{x}') + \sigma_n^2 I| - \frac{n}{2} \ln 2\pi$$

- Cannot guarantee a global optimum; try different initial conditions
- Constraint (hyper)parameters to some sensible limits
- **Note:** Choosing the right covariance function (e.g Squared Exponential, Matérn, Rational Quadratic, etc.) is not always easy and should be treated akin to a model selection problem

Learning the Covariance Parameters

- Let us define θ as a vector containing all (hyper)parameters
e.g $\theta = \{\alpha, l, \sigma_n^2\}$
- We choose θ that maximises the log marginal likelihood, that is:

$$\ln p(y|\mathbf{x}, \theta) = -\frac{1}{2}y^T(k(\mathbf{x}, \mathbf{x}') + \sigma_n^2 I)^{-1}y - \frac{1}{2} \ln |k(\mathbf{x}, \mathbf{x}') + \sigma_n^2 I| - \frac{n}{2} \ln 2\pi$$

- Cannot guarantee a global optimum; try different initial conditions
- Constraint (hyper)parameters to some sensible limits
- **Note:** Choosing the right covariance function (e.g Squared Exponential, Matérn, Rational Quadratic, etc.) is not always easy and should be treated akin to a model selection problem


Learning the Covariance Parameters

- Let us define θ as a vector containing all (hyper)parameters
e.g $\theta = \{\alpha, l, \sigma_n^2\}$
- We choose θ that maximises the log marginal likelihood, that is:


$$\ln p(y|\mathbf{x}, \theta) = -\frac{1}{2}y^T(k(\mathbf{x}, \mathbf{x}') + \sigma_n^2 I)^{-1}y - \frac{1}{2} \ln |k(\mathbf{x}, \mathbf{x}') + \sigma_n^2 I| - \frac{n}{2} \ln 2\pi$$

- Cannot guarantee a global optimum; try different initial conditions
- Constraint (hyper)parameters to some sensible limits
- **Note:** Choosing the right covariance function (e.g Squared Exponential, Matérn, Rational Quadratic, etc.) is not always easy and should be treated akin to a model selection problem


Modelling CO₂ Concentrations




Modelling CO₂ Concentrations




Modelling Gene Expression Time-Series




Modelling Gene Expression Time-Series



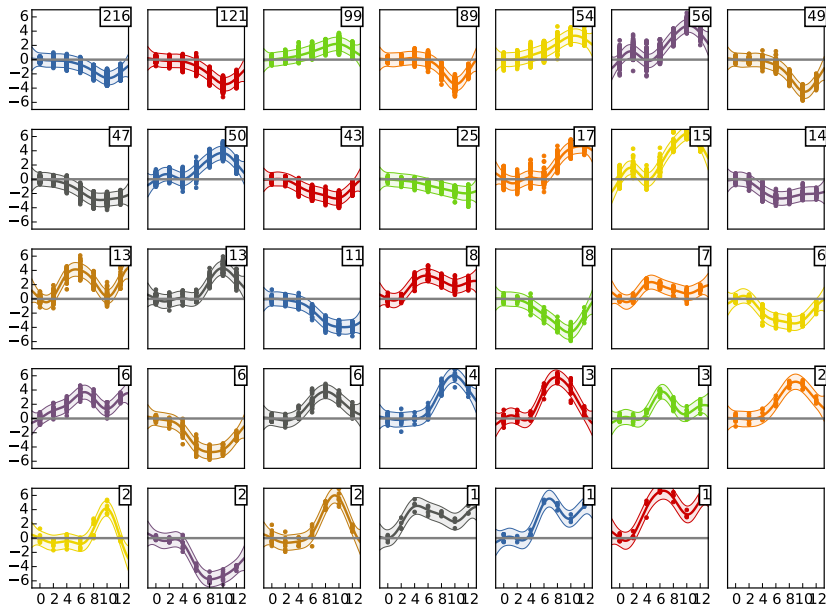
Modelling Gene Expression Time-Series



Ranking Gene Expression Time-Series




Clustering Gene Expression Time-Series



- Rasmussen and Williams book

Gaussian Processes for Machine Learning

Carl Edward Rasmussen and Christopher K. I. Williams
The MIT Press, 2006. ISBN 0-262-18253-X.



- <http://www.gaussianprocess.org/>
- The GPy Python Module by Neil Lawrence's Sheffield Group:
<https://github.com/SheffieldML/GPy>